



CS3281 / CS5281  
**Networking**

CS3281 / CS5281  
Spring 2024



Tel (615) 343-7472 | Fax (615) 343-7440  
1025 16th Avenue South Nashville, TN 37212  
[www.isis.vanderbilt.edu](http://www.isis.vanderbilt.edu)



# Intro to Sockets

- Sockets: method for IPC between applications
  - Can be on the same host
  - Can be on a different host connected by a network
- Typical organization: client-server
  - The client makes requests
    - Example: a web browser
  - The server responds to requests
    - Example: an Apache web server
- Communication involves a network protocol
  - Usually multiple layers of network protocols
- We'll cover TCP/IP
  - Also called the Internet protocol suite

# Big Picture: The Internet

- Began in 1960s: as network that could connect computers that were far away
  - Funding came from DARPA, and first ARPANET message was sent from UCLA to Stanford (350 miles) in 1969
- Originally linked research operations and CS departments
  - Spread to the commercial world in the 1990s and become “the Internet”
- Today: the Internet links millions of loosely connected, independent networks
- Data is through the networks in “packets” called IP (Internet Protocol) packets
  - Transported in one or more physical packets, like Ethernet or WiFi
  - Each IP packet passes through multiple gateways
    - Each gateway passes the packet to a gateway closer to the ultimate destination
- An internet (lowercase i) connects different computer networks
  - The Internet (capital I) refers to the TCP/IP internet that connects millions of computers
  - Some modern style guides do not capitalize “Internet.” We do here for conceptual clarity.

[https://en.wikipedia.org/wiki/Capitalization\\_of\\_Internet](https://en.wikipedia.org/wiki/Capitalization_of_Internet)




# The Internet (cont.)

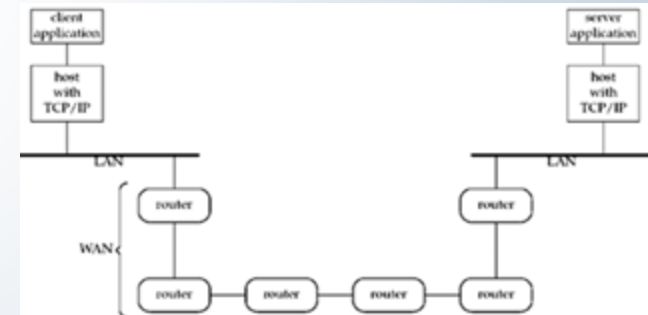
- The core protocol is the Internet Protocol
  - Defines a uniform transport mechanism and a common format for information in transit
  - IP packets are carried by different kinds of hardware using their own protocols
- The Transmission Control Protocol (TCP) sits on top of IP
  - TCP provides a reliable mechanism for sending arbitrarily long sequences of bytes
- Above TCP, higher-level protocols use TCP to provide services that we think of as “the Internet”
  - Examples: browsing, e-mail, file sharing
- All of these protocols taken together define the Internet

# Protocol Layers

- Example on the right:
  - Web servers and web clients communicate using TCP
  - TCP uses the Internet Protocol (IP)
  - IP uses a data link protocol (like Ethernet)
- The client and server use an application protocol
  - The transport layers use the TCP protocol
- Information flows down the protocol stack on one side, back up on the other
- Client and server are in user space
  - TCP, IP, data link in kernel space (usually)



On the same LAN



## On different LANs

# Sockets

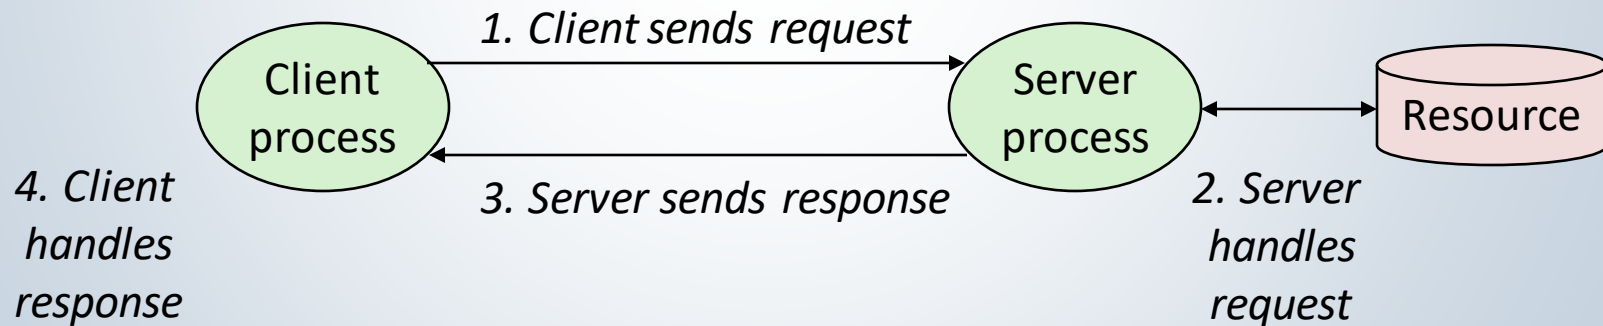
- What is a socket?
  - To the kernel, a socket is an endpoint of communication
  - To an application, a socket is a file descriptor that lets the application read/write from/to the network
    - **Remember:** All Unix I/O devices, including networks, are modeled as files
- Clients and servers communicate with each other by reading from and writing to socket descriptors



- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors

# A Client-Server Transaction

- Most network applications are based on the client-server model:
  - A **server** process and one or more **client** processes
  - Server manages some **resource**
  - Server provides **service** by manipulating resource for clients
  - Server activated by request from client (vending machine analogy)



*Note: clients and servers are processes running on hosts (can be the same or different hosts)*

# Sockets and Client/Server Communications

- Each application creates a socket
- The server binds its socket to a well-known address so clients can locate it

```
fd = socket(domain, type, protocol);
```

- Domain determines:
  - Format of address, and range of communication (same or different hosts)
    - AF\_UNIX, AF\_INET, AF\_INET6
- Type: stream or datagram
- Protocol: generally 0
  - Nonzero for some types like raw sockets (passes directly from data link to application)

Property	Socket type	
	Stream	Datagram
Reliable delivery?	Y	N
Message boundaries preserved?	N	Y
Connection-oriented?	Y	N



# Stream Sockets

- Stream sockets provide reliable, bidirectional, byte-stream communication
  - Reliable: Either the transmitted data arrives intact at the receiving end, or we receive notification of a probable failure in transmission
  - Bidirectional: data may be transmitted in either direction
  - Byte-stream: no message boundaries
    - Example: receiver doesn't know if the sender originally sent two 1-byte messages or one 2-byte message
- Operate in connected pairs (aka connection oriented)
  - *Peer* socket: socket at the other end of a connection
  - *Peer* address: address of that socket
  - *Peer* application: application using the peer socket
    - *Peer* is equivalent to *remote* or *foreign*

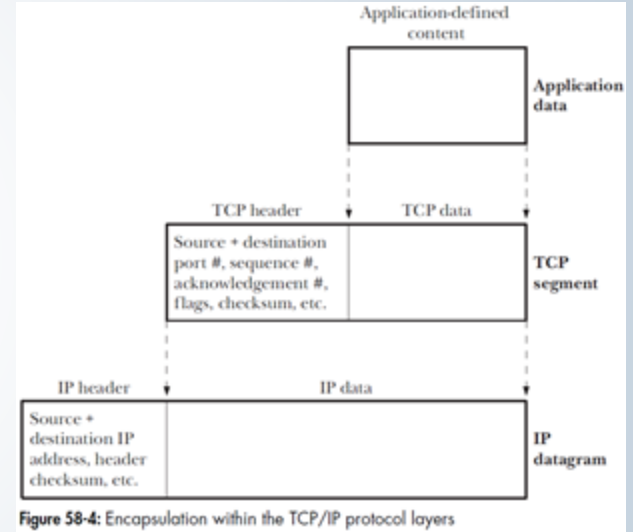
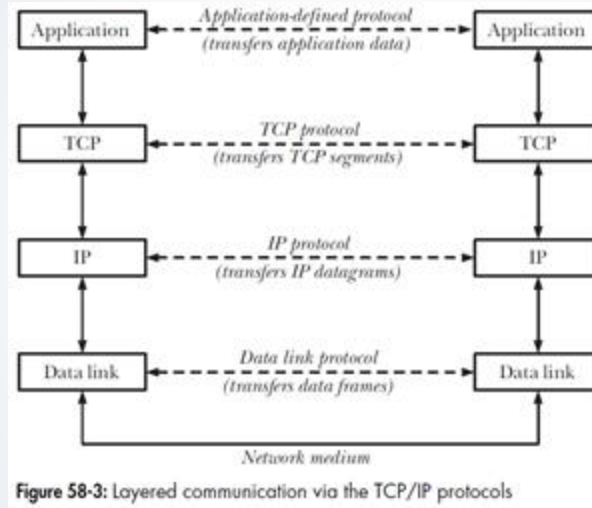

# Datagram Sockets

- Allow data to be exchanged in the form of messages called *datagrams*
  - Message boundaries are preserved
  - Data transmission is not reliable
    - Data may arrive out of order, be duplicated, or not arrive at all
  - Example of a connectionless socket
    - Doesn't need to be connected to another socket in order to be used
- In the Internet domain:
  - Datagram sockets use UDP
  - Stream sockets use\* TCP

*\*Almost always*



# Protocols and Communication



# Global IP Internet (Upper Case)


- Most famous example of an internet
- Based on the TCP/IP protocol family
  - IP (Internet Protocol) :
    - Provides *basic naming scheme* and unreliable *delivery capability* of packets (datagrams) from *host-to-host*
  - UDP (Unreliable Datagram Protocol)
    - Uses IP to provide *unreliable* datagram delivery from *process-to-process*
  - TCP (Transmission Control Protocol)
    - Uses IP to provide *reliable* byte streams from *process-to-process* over *connections*
- Accessed via a mix of Unix file I/O and functions from the *sockets interface*

# Computer Networks

- A *network* is a hierarchical system of boxes and wires organized by geographical proximity
  - SAN (System Area Network) spans cluster or machine room
    - Switched Ethernet, Quadrics QSW, ...
  - LAN (Local Area Network) spans a building or campus
    - Ethernet is most prominent example
  - WAN (Wide Area Network) spans country or world
    - Typically high-speed point-to-point phone lines
- An *internetwork* (*internet*) is an interconnected set of networks
  - The Global IP Internet (uppercase “I”) is the most famous example of an internet (lowercase “i”)
- Let’s see how an internet is built from the ground up


# Conceptual View of LANs

- For simplicity, hubs, bridges, and wires are often shown as a collection of hosts attached to a single wire:



# Next Level: Internets

- Multiple incompatible LANs can be physically connected by specialized computers called *routers*
- The connected networks are called an *internet* (lower case)



*LAN 1 and LAN 2 might be completely different, totally incompatible  
(e.g., Ethernet, Fibre Channel, 802.11\*, T1-links, DSL, ...)*

# The Notion of an Internet Protocol

- How is it possible to send bits across incompatible LANs and WANs?
- Solution: *protocol* software running on each host and router
  - Protocol is a set of rules that governs how hosts and routers should cooperate when they transfer data from network to network.
  - Smooths out the differences between the different networks



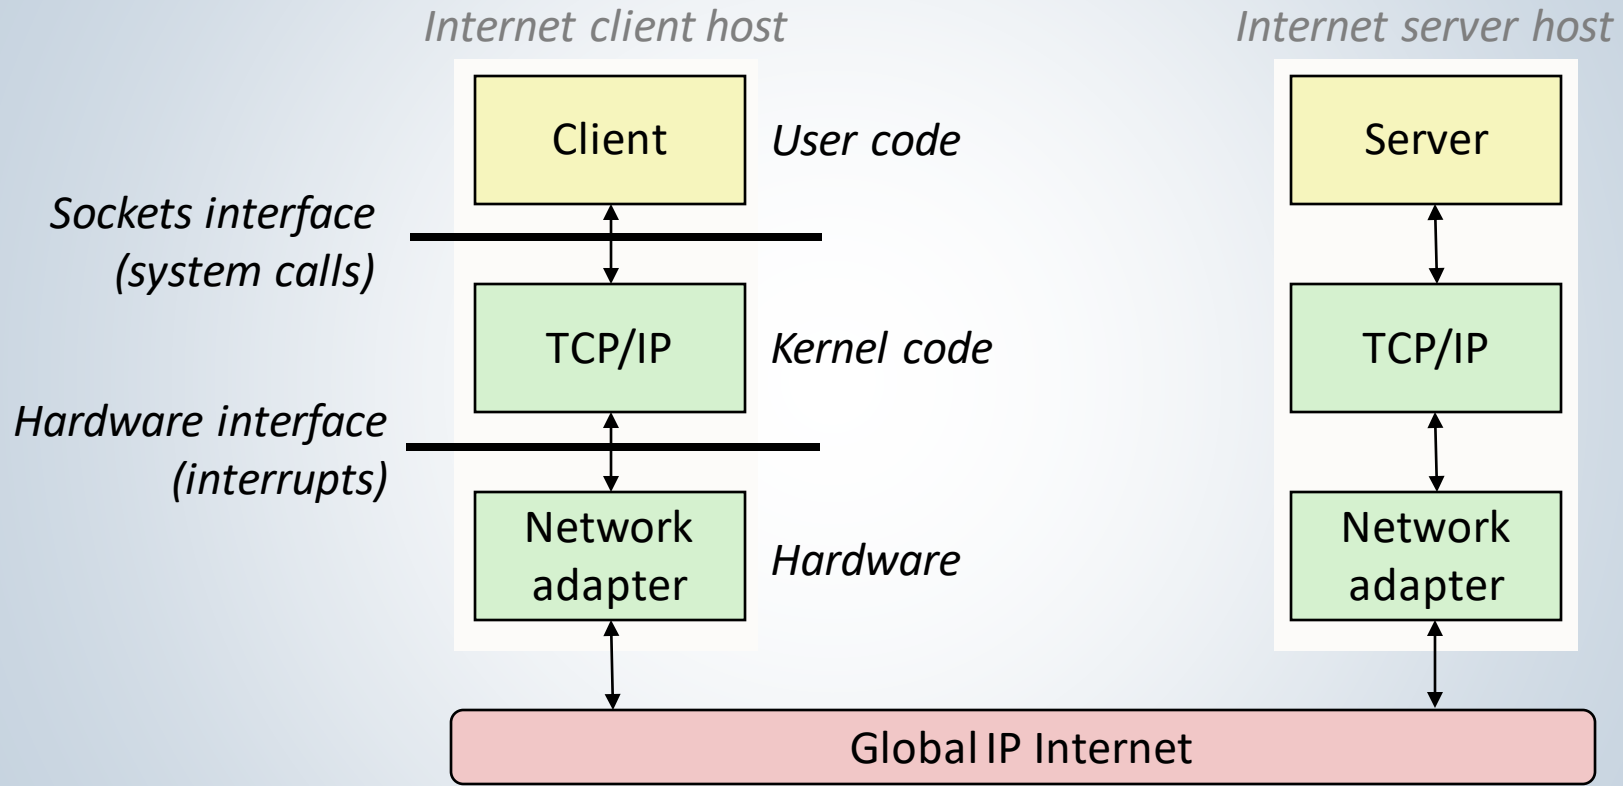
# What Does an Internet Protocol Do?

- Provides a *naming scheme*
  - An internet protocol defines a uniform format for **host addresses**
  - Each host (and router) is assigned at least one of these internet addresses that uniquely identifies it
- Provides a *delivery mechanism*
  - An internet protocol defines a standard transfer unit (**packet**)
  - Packet consists of **header** and **payload**
    - Header: contains info such as packet size, source and destination addresses
    - Payload: contains data bits sent from source host

# Other Issues

- We are glossing over a number of important questions:
  - What if different networks have different maximum frame sizes? (segmentation)
  - How do routers know where to forward frames?
  - How are routers informed when the network topology changes?
  - What if packets get lost?
- These (and other) questions are addressed by the area of systems known as *computer networking*

# Organization of an Internet Application



# A Programmer's View of the Internet

1. Hosts are mapped to a set of 32-bit *IP addresses*
  - 128.2.203.179
2. The set of IP addresses is mapped to a set of identifiers called Internet *domain names*
  - 129.59.107.38 is mapped to `www.isis.vanderbilt.edu`
3. A process on one Internet host can communicate with a process on another Internet host over a *connection*



# IP Addresses

- 32-bit IP addresses are stored in an *IP address struct*
  - IP addresses are always stored in memory in *network byte order* (big-endian byte order)
    - x86, ARM, risc-v all little endian
  - True in general for any integer transferred in a packet header from one machine to another
    - E.g., the port number used to identify an Internet connection

```
/* Internet address structure */  
struct in_addr {  
    uint32_t    s_addr; /* network byte order (big-endian) */  
};
```

# Dotted-Decimal Notation

- By convention, each byte in a 32-bit IP address is represented by its decimal value and separated by a period
  - IP address: 0x8002C2F2 = 128.2.194.242



# Internet Connections

- Clients and servers communicate by sending streams of bytes over **connections**. Each connection is:
  - *Point-to-point*: connects a pair of processes.
  - *Full-duplex*: data can flow in both directions at the same time,
  - *Reliable*: stream of bytes sent by the source is eventually received by the destination in the same order it was sent.
- A **socket** is an endpoint of a connection
  - *Socket address* is an **IPAddress:port** pair
- A **port** is a 16-bit integer that identifies a process:
  - **Ephemeral port**: Assigned automatically by client kernel when client makes a connection request.
  - **Well-known port**: Associated with some **service** provided by a server (e.g., port 80 is associated with Web servers)

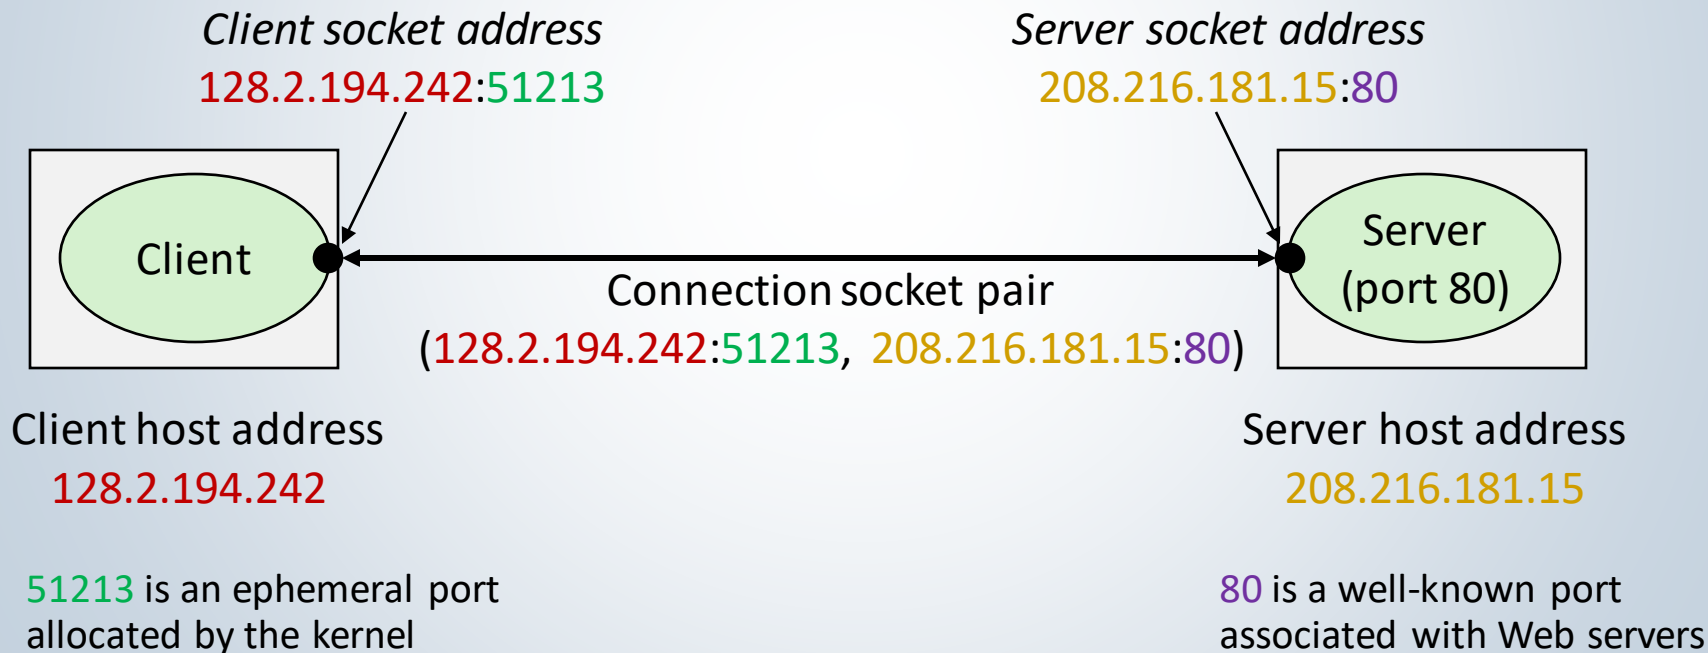
# Well-Known Ports and Service Names

- Popular services have permanently assigned *well-known ports* and corresponding *well-known service names*:
  - echo server: 7/echo
  - ssh servers: 22/ssh
  - email server: 25/smtp
  - Web servers: 80/http, 443/https
- Mappings between well-known ports and service names is contained in the file `/etc/services` on each Linux machine.





# Anatomy of a Connection

- A connection is uniquely identified by the socket addresses of its endpoints (*socket pair*)
  - (cliaddr:cliport, servaddr:servport)



# Using Ports to Identify Services



# Sockets Interface

- Set of system-level functions used in conjunction with Unix I/O to build network applications.
- Created in the early 80's as part of the original Berkeley distribution of Unix that contained an early version of the Internet protocols.
- Available on all modern systems
  - Unix variants, Windows, OS X, IOS, Android, ARM

# Socket-Address Structures

- Generic socket address:
  - For address arguments to **connect**, **bind**, and **accept**
  - Necessary only because C did not have generic (**void \***) pointers when the sockets interface was designed
  - For casting convenience, we adopt the Stevens convention:

**typedef struct sockaddr SA;**

```
struct sockaddr {  
    uint16_t    sa_family;    /* Protocol family */  
    char        sa_data[14]; /* Address data. */  
};
```

sa\_family

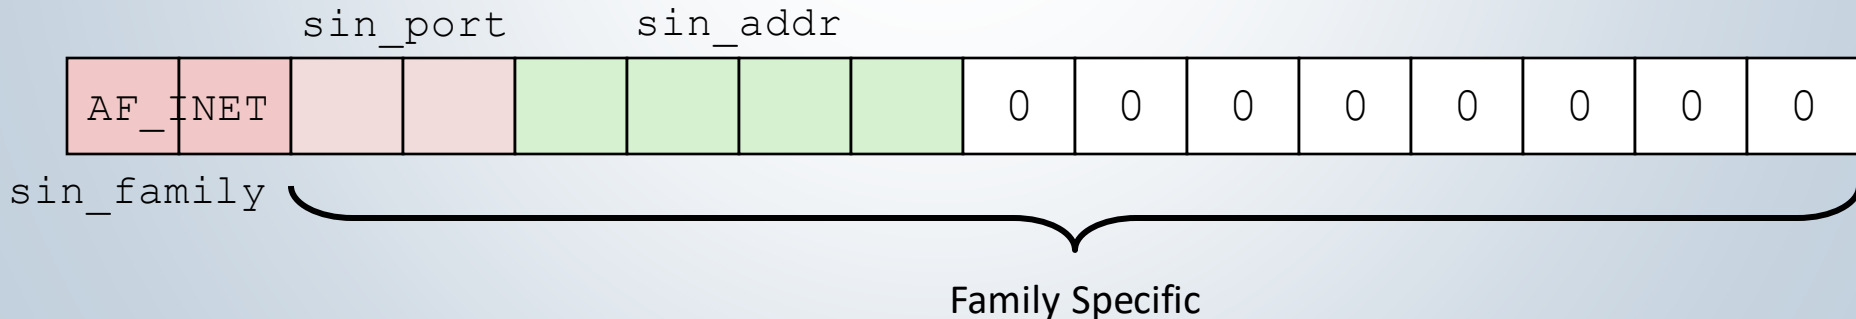


Family Specific

# Socket Address Structures

- Internet-specific socket address:
  - Must cast (`struct sockaddr_in *`) to (`struct sockaddr *`) for functions that take socket address arguments.

```
struct sockaddr_in {  
    uint16_t    sin_family; /* Protocol family (always AF_INET) */  
    uint16_t    sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr; /* IP addr in network byte order */  
    unsigned char sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```



# Key Socket Calls

- `socket()` creates a new socket
- `bind()` binds a socket to an address
- `listen()` lets a TCP socket to accept incoming connections from other sockets
- `accept()` accepts a connection from a peer application
- `connect()` establishes a connection with another socket
- Socket I/O can be done using
  - `read()` and `write()`, or
  - `send()`, `recv()`, `sendto()`, `recvfrom()`

# Key Socket Calls

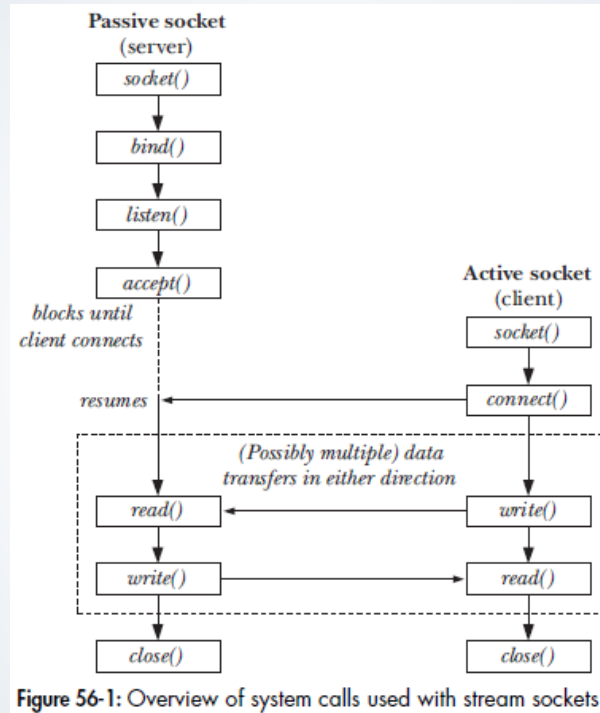


Figure 56-1: Overview of system calls used with stream sockets

Figure from *The Linux Programming Interface* by Michael Kerrisk

# Client-Server Example

```
void server()
{
    int fd;
    struct sockaddr_in in_addr;
    memset(&in_addr, 0, sizeof(struct sockaddr_in));

    if ((fd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
        exit_with_error("Server error with socket");
    }

    in_addr.sin_family = AF_INET;
    in_addr.sin_port = 5001;
    inet_pton(AF_INET, "0.0.0.0", &in_addr.sin_addr);

    if (bind(fd, (struct sockaddr *) &in_addr, sizeof(struct sockaddr_in))) {
        exit_with_error("Server error with bind");
    }

    if (listen(fd, 0)) {
        exit_with_error("Server error with listen");
    }

    int port;
    struct sockaddr_in client_info;
    socklen_t client_size = sizeof(client_info);
    memset(&client_info, 0, sizeof(client_info));

    if ((port = accept(fd, (struct sockaddr *) &client_info, &client_size)) == -1) {
        exit_with_error("Server error with accept");
    }

    int ret, total = 0;
    char buf[100];
    while ((ret = read(port, buf + total, sizeof(buf) - 1 - total)) > 0) {
        if (ret == -1) {
            exit_with_error("Server error with read");
        }
        total += ret;
    }

    buf[total] = 0;
    printf("Server received: %s, total = %d\n", buf, total);
}
```

```
void client(int port)
{
    int fd;
    struct sockaddr_in in_addr;
    memset(&in_addr, 0, sizeof(struct sockaddr_in));

    if ((fd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
        exit_with_error(0);
    }

    in_addr.sin_family = AF_INET;
    in_addr.sin_port = 5001;
    inet_pton(AF_INET, "127.0.0.1", &in_addr.sin_addr);

    if (connect(fd, (struct sockaddr *) &in_addr, sizeof(struct sockaddr_in))) {
        exit_with_error("Client error with connect");
    }

    int ret, sent = 0;
    char *msg = "Hello, server!";
    while (sent != strlen(msg)) {
        ret = write(fd, msg + sent, strlen(msg) - sent);
        if (ret == -1) {
            exit_with_error("Client error with write");
        }
        else {
            sent += ret;
        }
        printf("sent = %d\n", sent);
    }

    void exit_with_error(char *msg)
    {
        perror(msg);
        exit(1);
    }
}
```



# Client-Server Example

## IPv4 socket address structure

```
struct in_addr {
    in_addr_t s_addr;
};

struct sockaddr_in {
    sa_family_t sin_family;
    in_port_t sin_port;
    struct in_addr sin_addr;
    unsigned char __pad[X];
};
```

/\* IPv4 4-byte address \*/  
/\* Unsigned 32-bit integer \*/  
/\* IPv4 socket address \*/  
/\* Address family (AF\_INET) \*/  
/\* Port number \*/  
/\* IPv4 address \*/  
/\* Pad to size of 'sockaddr' structure (16 bytes) \*/

## Generic socket address structure

```
struct sockaddr {
    sa_family_t sa_family;
    char sa_data[14];
};
```

/\* Address family (AF\_\* constant) \*/  
/\* Socket address (size varies according to socket domain) \*/

```
void server()
{
    int fd;
    struct sockaddr_in in_addr;
    memset(&in_addr, 0, sizeof(struct sockaddr_in));

    if ((fd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
        exit_with_error("Server error with socket");
    }

    in_addr.sin_family = AF_INET;
    in_addr.sin_port = 5001;
    inet_pton(AF_INET, "0.0.0.0", &in_addr.sin_addr);

    if (bind(fd, (struct sockaddr *) &in_addr, sizeof(struct sockaddr_in))) {
        exit_with_error("Server error with bind");
    }

    if (listen(fd, 0)) {
        exit_with_error("Server error with listen");
    }

    int port;
    struct sockaddr_in client_info;
    socklen_t client_size = sizeof(client_info);
    memset(&client_info, 0, sizeof(client_info));

    if ((port = accept(fd, (struct sockaddr *) &client_info, &client_size)) == -1) {
        exit_with_error("Server error with accept");
    }

    int ret, total = 0;
    char buf[100];
    while ((ret = read(port, buf + total, sizeof(buf) - 1 - total)) > 0) {
        if (ret == -1) {
            exit_with_error("Server error with read");
        }
        total += ret;
    }

    buf[total] = 0;
    printf("Server received: %s, total = %d\n", buf, total);
}
```

```
#include <sys/socket.h>
```

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

Returns 0 on success, or -1 on error

```
#include <sys/socket.h>
```

```
int listen(int sockfd, int backlog);
```

Returns 0 on success, or -1 on error

```
#define __ss_aligntype uint32_t
struct sockaddr_storage {
    sa_family_t ss_family;
    __ss_aligntype __ss_align;
    char __ss_padding[SS_PADSIZE];
};
```

/\* On 32-bit architectures \*/  
/\* Force alignment \*/  
/\* Pad to 128 bytes \*/

Large enough for IPv4 or IPv6

```
#include <sys/socket.h>
```

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

Returns file descriptor on success, or -1 on error

# Client-Server Example

```
void client(int port)
{
    int fd;
    struct sockaddr_in in_addr;
    memset(&in_addr, 0, sizeof(struct sockaddr_in));

    if ((fd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
        exit_with_error(0);
    }

    in_addr.sin_family = AF_INET;
    in_addr.sin_port = 5001;
    inet_pton(AF_INET, "127.0.0.1", &in_addr.sin_addr);

    if (connect(fd, (struct sockaddr *)&in_addr, sizeof(struct sockaddr_in))) {
        exit_with_error("Client error with connect");
    }

    int ret, sent = 0;
    char *msg = "Hello, server!";
    while (sent != strlen(msg)) {
        ret = write(fd, msg + sent, strlen(msg) - sent);
        if (ret == -1)
            exit_with_error("Client error with write");
        else
            sent += ret;
        printf("sent = %d\n", sent);
    }
}
```

```
void exit_with_error(char *msg)
{
    perror(msg);
    exit(1);
}
```

## IPv4 socket address structure

```
struct in_addr {
    in_addr_t s_addr;
};

struct sockaddr_in {
    sa_family_t sin_family;
    in_port_t sin_port;
    struct in_addr sin_addr;
    unsigned char __pad[X];
};
```

#include <arpa/inet.h>

int inet\_pton(int domain, const char \*src\_str, void \*addrptr);

Returns 1 on successful conversion, 0 if *src\_str* is not in presentation format, or -1 on error

#include <sys/socket.h>

int connect(int sockfd, const struct sockaddr \*addr, socklen\_t addrlen);

Returns 0 on success, or -1 on error

Basic write() library call uses a generic file descriptor

Print the last error encountered during a system call or library function