



CS3281 / CS5281 I/O Devices

CS3281 / CS5281
Spring 2024

**Some lecture slides borrowed and adapted from
Andrea Arpaci-Dusseau*



Tel (615) 343-7472 | Fax (615) 343-7440
1025 16th Avenue South Nashville, TN 37212
www.isis.vanderbilt.edu



Motivation

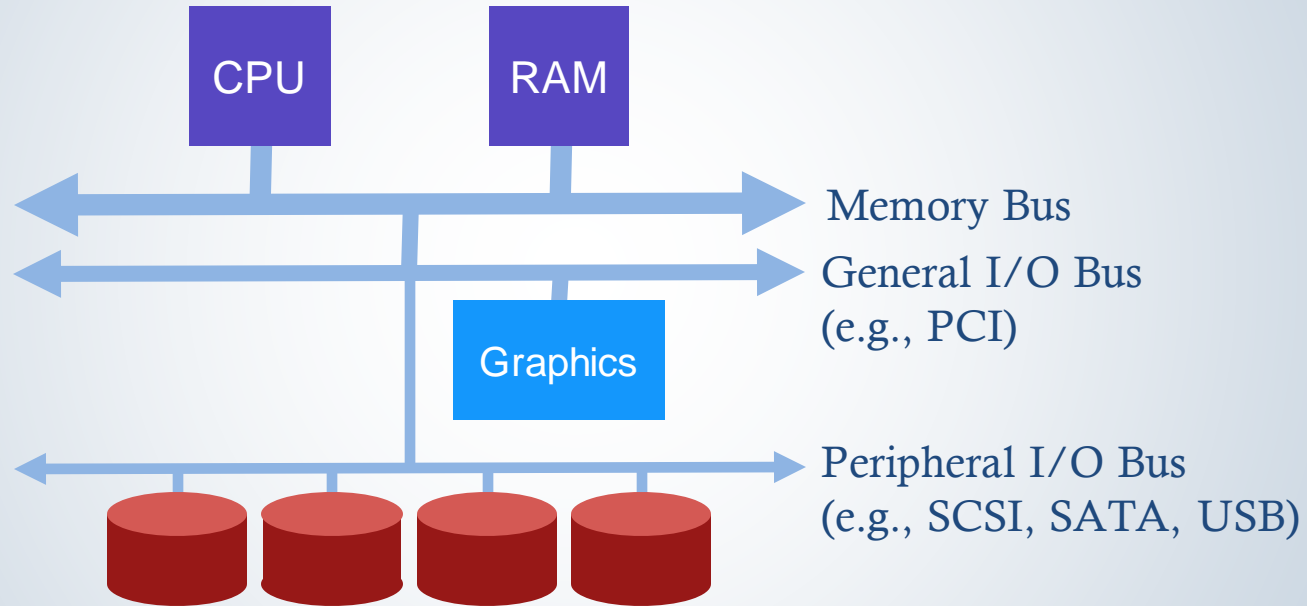
What good is a computer without any I/O devices?

- e.g., keyboard, display, disks

We want:

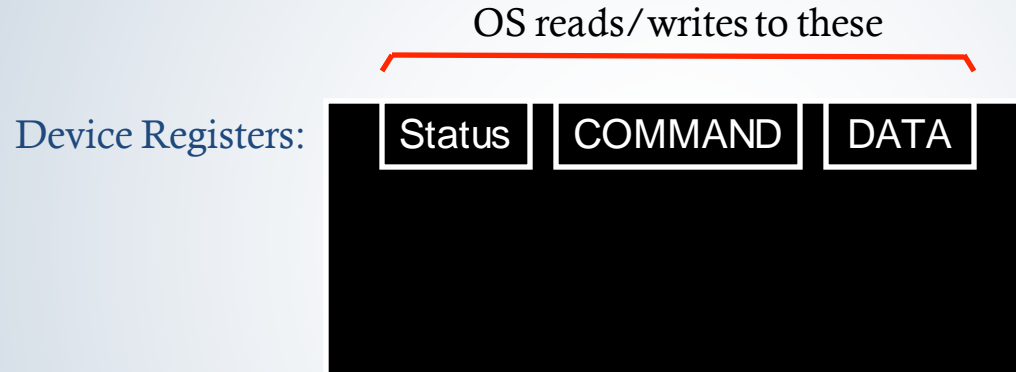
- **H/W** that will let us plug in different devices
- **OS** that can interact with different combinations
- I/O also allows for *persistence*
 - RAM is *volatile*, i.e., contents are lost when the machine restarts

Hardware support for I/O

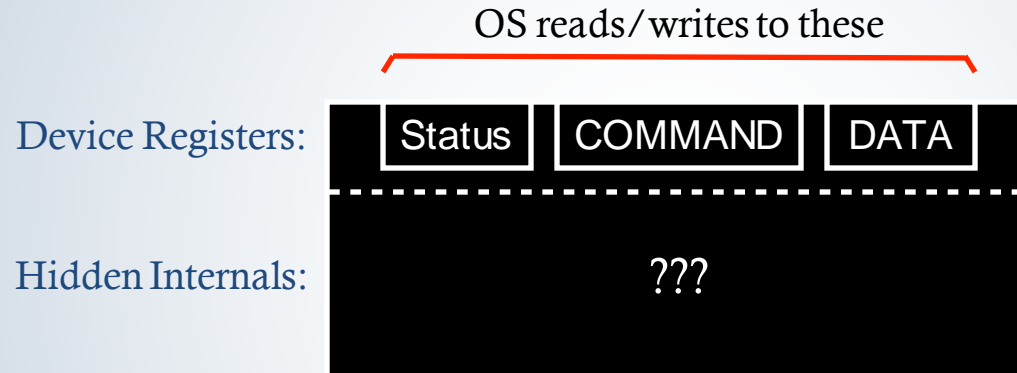


Why use hierarchical buses?

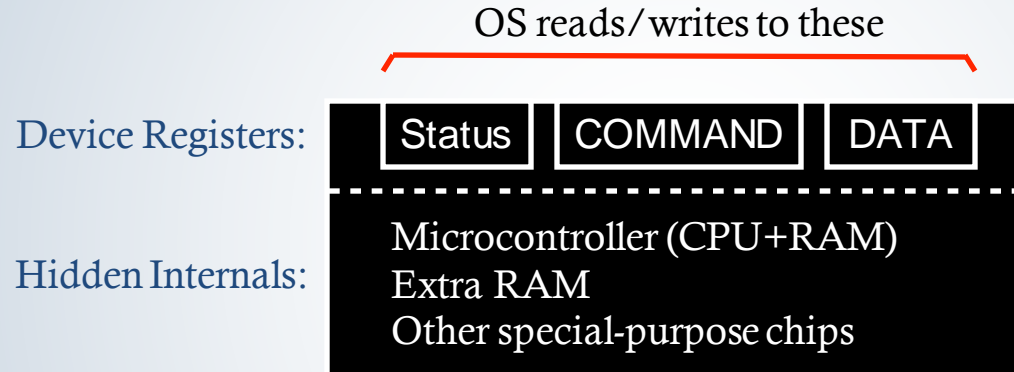
Canonical Device



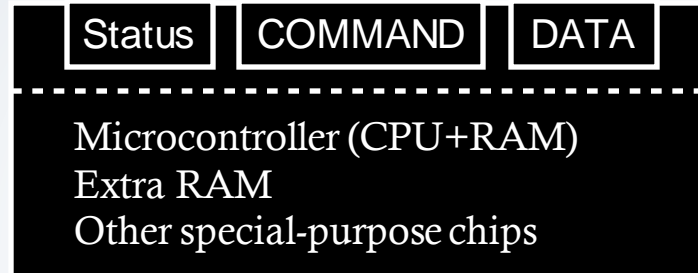
Canonical Device



Canonical Device



Example Write Protocol



```
while (STATUS == BUSY)
    ; // spin
```

Write data to DATA register

Write command to COMMAND register

```
while (STATUS == BUSY)
    ; // spin
```

This is called polling when the processor “asks” what the hardware is doing, often continuously

CPU:

Disk:

```
while (STATUS == BUSY)      // 1
```

```
;
```

```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
;
```


CPU: 

Disk: 

```
while (STATUS == BUSY)      // 1
```

```
;
```

```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
;
```

A wants to do I/O



CPU: A

Disk: C

```
while (STATUS == BUSY)      // 1
```

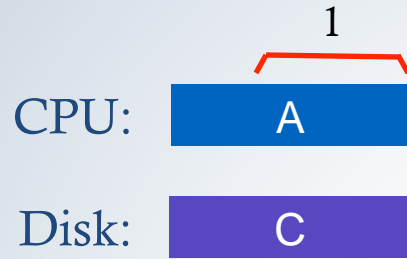
```
;
```

```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
;
```



```
while (STATUS == BUSY)    // 1
```

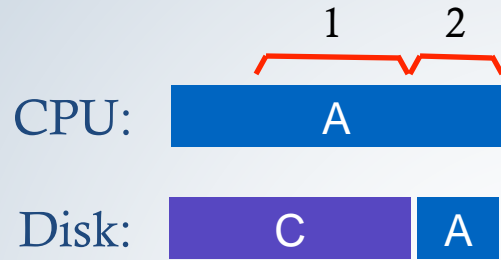
```
;
```

```
Write data to DATA register    // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)    // 4
```

```
;
```



```
while (STATUS == BUSY)      // 1
```

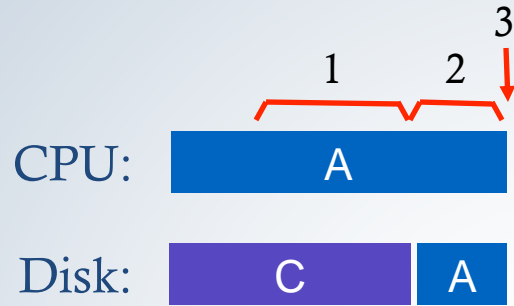
```
;
```

```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
;
```



```
while (STATUS == BUSY)      // 1
```

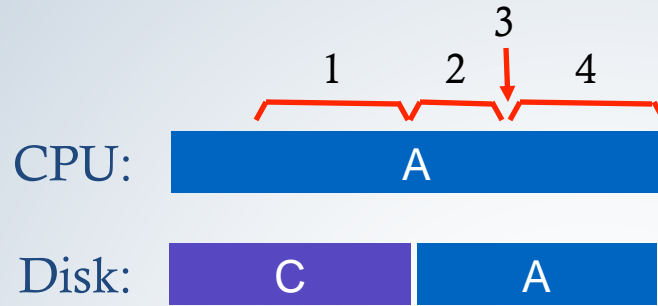
```
;
```

```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
;
```



```
while (STATUS == BUSY)      // 1
```

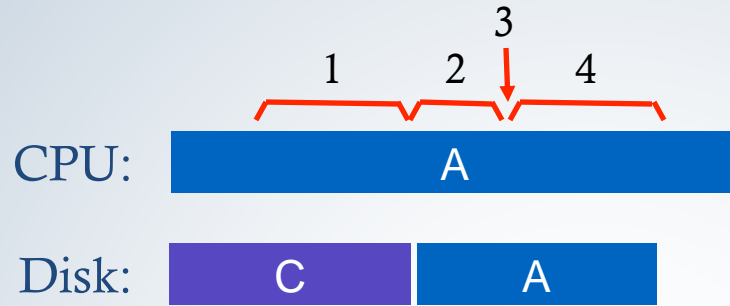
```
;
```

```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
;
```



```
while (STATUS == BUSY)      // 1
```

```
;
```

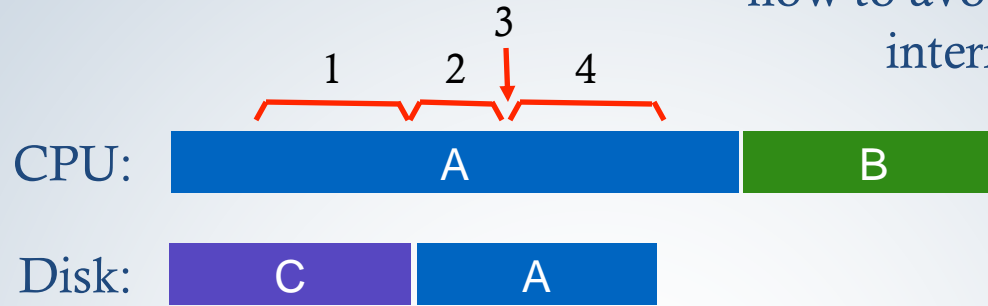
```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
;
```

how to avoid spinning?
interrupts!



```
while (STATUS == BUSY)      // 1
```

```
    wait for interrupt;
```

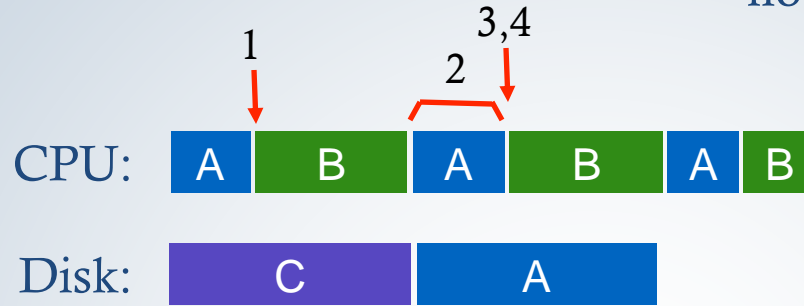
```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
    wait for interrupt;
```


how to avoid spinning?
interrupts!



```
while (STATUS == BUSY)    // 1
```

```
    wait for interrupt;
```

```
Write data to DATA register    // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)    // 4
```

```
    wait for interrupt;
```

Interrupts vs. Polling

Are interrupts ever worse than polling?

Fast device: Better to spin than take interrupt and context-switching overhead

- Device time unknown? Hybrid approach (spin then use interrupts)

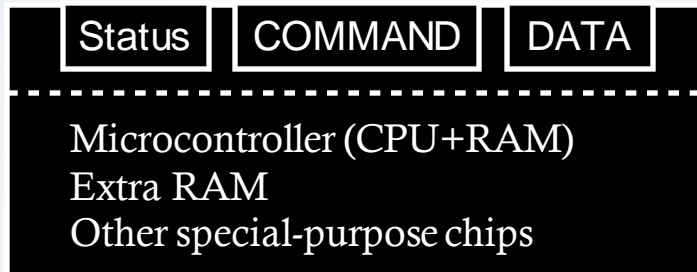
Flood of interrupts arrive

- Can lead to livelock (always handling interrupts)
- Better to ignore interrupts while making some progress handling them

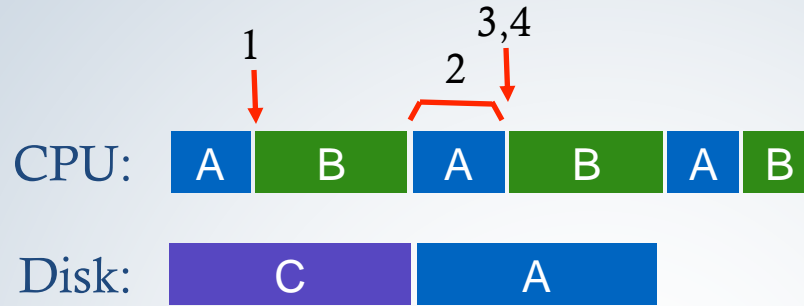
Other improvement

- Interrupt coalescing (hardware batches together several interrupts)

Protocol Variants



- **Status checks:** polling vs. interrupts
- **Data:** programmed I/O (PIO) vs. direct memory access (DMA)
- **Control:** special instructions vs. memory-mapped I/O



```
while (STATUS == BUSY)      // 1
```

```
    wait for interrupt;
```

```
    Write data to DATA register    // 2
```

```
    Write command to COMMAND register // 3
```

```
    while (STATUS == BUSY)      // 4
```

```
        wait for interrupt;
```

what else can we optimize?

data transfer!

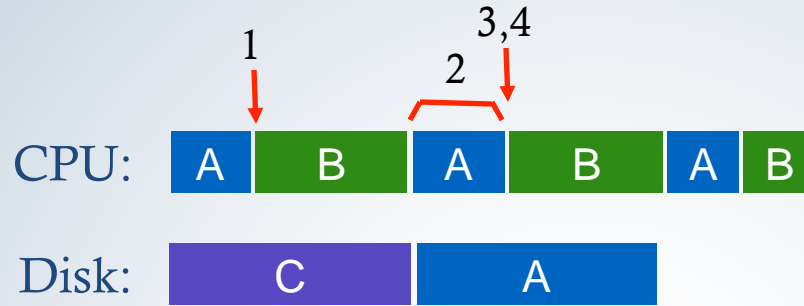
Programmed I/O vs. Direct Memory Access

PIO (Programmed I/O):

- CPU directly tells device what the data is
- One instruction for each byte/word
- Efficient for a few bytes/words, but *scales terribly*

DMA (Direct Memory Access):

- CPU leaves data in memory
- Device reads/writes data directly from/to memory
- One instruction to send a pointer to the data to send
- Efficient for large data transfers



```
while (STATUS == BUSY)      // 1
```

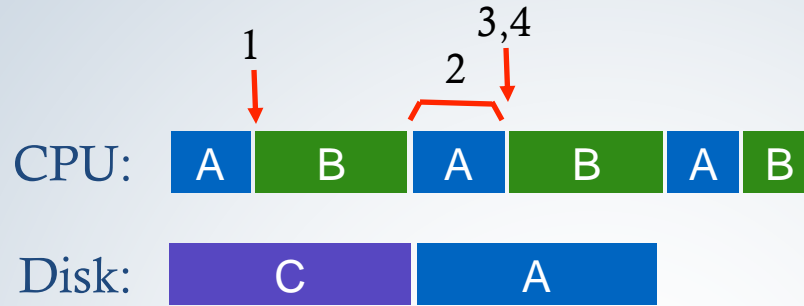
```
    wait for interrupt;
```

```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
    wait for interrupt;
```



```
while (STATUS == BUSY)      // 1
```

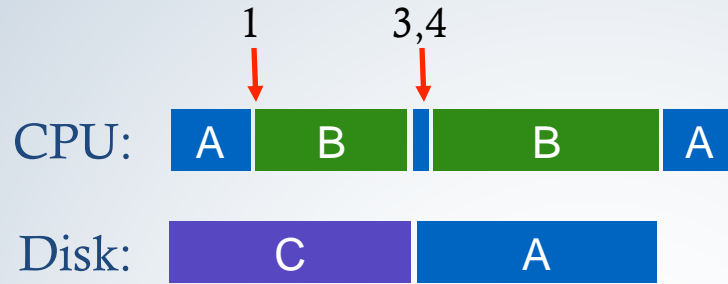
```
    wait for interrupt;
```

```
    Write data to DATA register // 2
```

```
    Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
    wait for interrupt;
```



```
while (STATUS == BUSY)      // 1
```

```
    wait for interrupt;
```

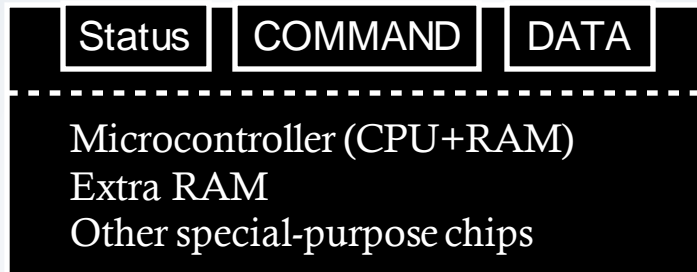
```
    Write data to DATA register // 2
```

```
    Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
    wait for interrupt;
```

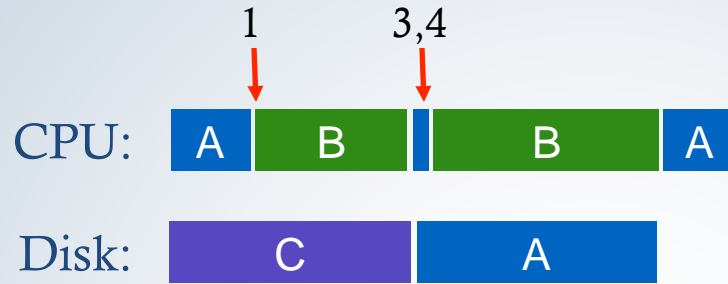

Protocol Variants



Status checks: polling vs. interrupts

Data: PIO vs. DMA

Control: special instructions vs. memory-mapped I/O



```
while (STATUS == BUSY)      // 1
```

```
    wait for interrupt;
```

```
Write data to DATA register // 2
```

```
Write command to COMMAND register // 3
```

```
while (STATUS == BUSY)      // 4
```

```
    wait for interrupt;
```

how does OS read and write registers?

Special Instructions vs. Mem-Mapped I/O

Special instructions

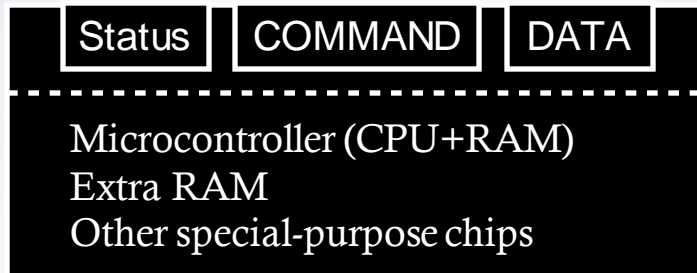
- each device has a port
- in/out instructions (x86) communicate with device

Memory-Mapped I/O

- H/W maps registers into address space
- loads/stores sent to device

Doesn't matter much (both are used)

Protocol Variants



Status checks: polling vs. interrupts

Data: PIO vs. DMA

Control: special instructions vs. memory-mapped I/O

Variety is a Challenge

Problem:

- many, many devices
- each has its own protocol

How can we avoid writing a slightly different OS for each H/W combination?

Write device driver for each device

Drivers are **70%** of Linux source code